

Mapping Mining Areas in the Tropics from 2016–2024

Philipp Sepin¹, Lukas Vashold¹, and Nikolas Kuschnig^{1,2,*}

¹Vienna University of Economics and Business

²Monash University

*Corresponding author: Nikolas Kuschnig (nikolas.kuschnig@monash.edu)

ABSTRACT

Mining provides crucial materials for the global economy and the energy transition but can have severe environmental and social impacts. Available analyses of these impacts are limited by a lack of data on mining activity, particularly in the regions most affected. In this paper, we present a detailed panel dataset mapping mining sites along the tropical belt from 2016 to 2024. Our approach uses a machine learning model, trained on over 25,000 mining polygons from the literature, to automatically segment mining areas in high-resolution (< 5 m) satellite imagery. The dataset maps over 145,000 mining polygons covering an average area of 65,000 km² annually, with an accuracy of 87.7% and precision of 84.1%. Our approach allows for accurate, precise, and consistent delineation, and can be scaled to other locations and periods. The dataset enables detailed analyses of local environmental, social, and economic impacts of mining in regions where conventional data are scarce or incomplete.

Transition minerals play a crucial role in climate action and are necessary for the switch towards cleaner production, storage, and distribution of energy.¹ Globally, mining operations are expanding to meet the growing demand for raw materials, often encroaching upon vulnerable regions.² Projections suggest that this expansion will accelerate drastically in pursuit of the Paris Agreement and subsequent climate conferences.³ Yearly extraction of critical minerals is projected to increase by 150–450% depending on the mineral, with a cumulative total of material extracted reaching 1.8–3.5 billion tons by 2050.⁴

Understanding the impacts of increasing mineral extraction is crucial but relies on comprehensive data that is often lacking. On the one hand, mining is linked to several adverse environmental and social effects, including deforestation, loss of biodiversity, soil erosion, water pollution, air contamination, corruption, and violent conflicts.^{5–14} On the other hand, mining can present economic opportunities, having been shown to increase wealth levels, asset ownership, and incomes as well as related socioeconomic indicators,^{15–18} although the effects on local and regional development are ambiguous^{19–22} and may vary by mineral type and target market.²³ Yet, if managed successfully, the demand for minerals could positively affect economic development and help facilitate the delivery of the sustainable development goals (SDGs).²⁴

The effective management of mining impacts necessitates detailed information on the location, areal extent, and activity over time. Despite previous efforts,^{25–29} information on the status and especially the development of mining over time remains scarce³⁰ and comprehensive analyses are impeded.

In this paper, we introduce a panel dataset of mining areas within the tropical belt. The dataset spans from 2016 to 2024 and consists of mining polygons that are automatically delineated using state-of-the-art machine learning (ML) methods. Our methodology employs a transformer-based segmentation model,³¹ trained on an extensive dataset of mining polygons from the existing literature.^{26,28} We apply this model to delineate polygons of known mine sites using frequent, high-resolution satellite imagery from Planet, which is provided under Norway's International Climate and Forest Initiative (NICFI).

This approach provides a comprehensive dataset of mining polygons, which tracks yearly footprint

changes of previously identified mining sites. This data enables large-scale analyses of the various impacts of mineral extraction — particularly in regions where such data has historically been scarce. The nature of our approach allows for consistent and accurate predictions that can be easily extended across different locations and temporal ranges.

Results and Discussion

Our dataset covers mining sites in the tropical belt and tracks them at annual frequency from 2016 until 2024. Mining sites and their development over time are captured by a total of over 147,000 mining polygons that cover an average area of 66,400 km² over the observed period. [Figure 1](#) illustrates the spatial coverage of the dataset and presents the change in the area of monitored mining sites.

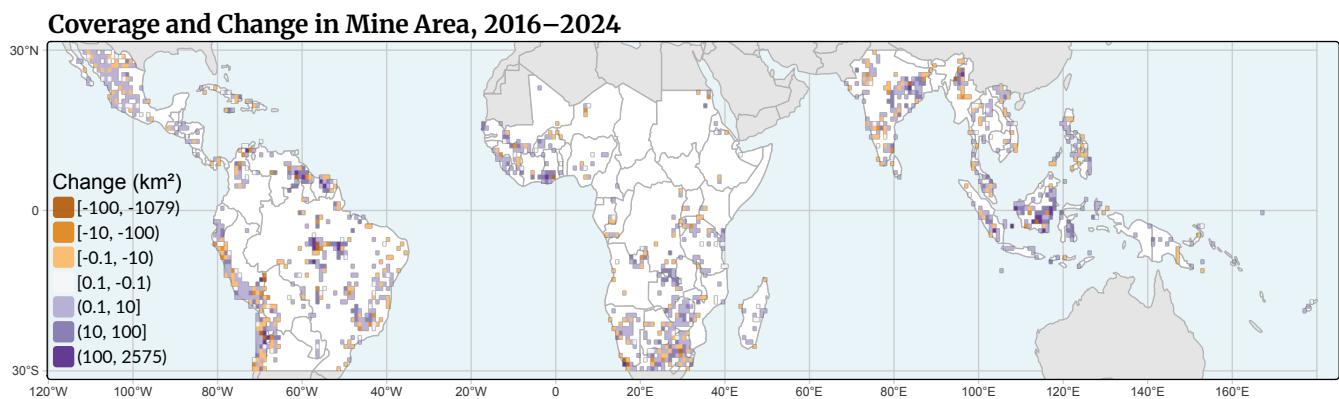


Figure 1. Change in delineated mining areas over 2016–2024. Colored regions indicate the coverage of our dataset (i.e., the tropical belt), while (non-white) pixels indicate the change in the area of tracked mining sites (km²).

In this section, we describe insights that can be obtained from our dataset as well as benefits from our modeling approach and discuss their usage, limitations, and areas for future work. We begin by focusing on the temporal dimension of our dataset, which allows researchers to track mining sites and their impacts over time. We illustrate how this information can be used to monitor the activity, expansion, and — potentially — renaturalization of individual mining sites. Afterward, we explore the development of mining areas at aggregated levels. Then, we focus on our modelling approach. Segmentation models can distill a consensus from heterogeneous training data, generating consistent polygons, and lower the cost of delineating polygons considerably. Finally, we discuss the limitations of this study and the introduced dataset, as well as directions for future research.

Tracking a mine site

[Figure 2](#) shows the development of the Toka Tindung mining site from 2016 to 2024 based on annual satellite imagery as well as the corresponding annual model delineations. The mining site is located in the Indonesian province of North Sulawesi, approximately 35 kilometers east of the provincial capital, Manado. Commercial production began in 2011, and production and processing capacity have been added since, making Toka Tindung one of the largest gold mining operations in Southeast Asia today. Over the period from 2016–2024, multiple upgrades of the processing plant more than quadrupled the project's throughput capacity to four million tons per year.³²

The increase in capacity is clearly reflected in [Figure 2](#) and captured by our predictions. The expansion of the main pits of Toka Tindung has been accompanied by the addition of infrastructure, such as water




Figure 2. Annual predictions (2016–2024) of the extent of the Toka Tindung gold mining project in Indonesia ($1^{\circ}35'N$ $125^{\circ}06'E$) are plotted over the corresponding annual satellite image (Planet/NICFI). Imagery ©2016–2024 Planet Labs Inc.

storage facilities. Over 2016–2024, the Toka pit in the north and the Kopra, Blambangan, and Araren pits in the south have expanded and grown closer to each other and expanded considerably. Later developments were concentrated in the southern parts of the mine site, including the Araren pit in the bottom-right of the images.

In the existing mining polygon datasets that we use to train our model,^{26,28} none of these developments are reflected (see [Figure S1](#) of the Supplementary Information for a visualization). The southern pits are missing from one of the sources,²⁶ indicating the use of satellite imagery from before 2019. The other source²⁸ delineates these elements individually, resulting in five polygons. By contrast, our predictions generally result in one large polygon, stemming from the broad nature of the ground truth used for training. Overall, our mining polygons consistently trace the expansion of Toka Tindung, better reflecting its actual state in a given year.

Changes of mining sites over time

Mining is prevalent throughout the tropics but is well known to be highly clustered in a few areas.^{26,28} Similarly, it is well known that mineral extraction is booming,⁴ but less is known about the changing footprint of mining areas. Below, we first show how our dataset reveals substantial changes in the location and extent of mining clusters in the period from 2016–2024. Then, we use the dataset to assess the development of mining sites over time.

In Latin America, Brazil, and Chile feature the largest and rapidly growing areal footprints of mining sites (see [Table 1](#)). As visible in [Figure 1](#), the Amazon rainforest has been a particular hotspot for increasing footprints, especially along the Brazilian ‘arc of deforestation’ at the southern and eastern edges of the Amazon and its northern edges, including the border region between Guyana, Venezuela, and Suriname. Traditional mining regions along the Chilean Andes show alternating patches of increases and decreases in the size of mining sites. Other countries in the Americas, such as Colombia and Mexico, show considerable, but more dispersed, growth in the footprint of mining sites.

Country	Mining Area (km ²)								
	13,266	12,425	12,376	13,761	15,281	16,859	15,746	17,497	17,801
Indonesia	13,266	12,425	12,376	13,761	15,281	16,859	15,746	17,497	17,801
Brazil	7,829	8,785	8,002	8,023	8,548	8,215	9,764	9,342	9,360
Chile	6,293	5,748	5,746	7,161	7,269	6,434	7,561	7,176	7,218
South Africa	7,469	6,624	5,901	6,706	6,324	6,044	6,434	6,459	6,558
India	3,502	3,800	3,590	4,140	4,136	3,959	4,322	4,367	4,114

Table 1. Mining polygon area in the five countries with the largest mining footprint in the dataset over time. A complete version with all countries with a footprint of over 100 km² in the dataset is available as [Table S1](#) in the Supplementary Information.

In the western parts of Africa, pronounced growth of areal footprints can be observed in traditional gold mining areas. Ghana is a primary hotspot, but mining sites in eastern Guinea, Sierra Leone, and Burkina Faso have also expanded consistently throughout the observation period. In the south, mining sites have expanded overall, but there has been a slight shift from sites in South Africa, where the overall mining polygon area has been stagnant, to Namibia, Zimbabwe, and Botswana. The Congo River basin, where mining operations have expanded rapidly,³³ is only partially captured by our dataset. Mining sites in the south are covered and growing strongly, while the often smaller and informal mining sites in the rest of the country are poorly covered by our sources and, hence, dataset.

In Asia, the islands of Borneo and Sulawesi in Indonesia, eastern India, and northern Myanmar are particular hotspots for growing mining sites. Other areas such as northern Thailand, Laos, and the Philippines also show expansion of mining areas, albeit at a more moderate pace. In contrast, parts of the Malay Peninsula, southern Myanmar, and southern India show a declining aerial extent of mining operations over the observation period.

Mineral extraction is expanding globally⁴ and Figure 3 shows how this coincides with increasing total area and number of mining polygons in our covered locations. We focus on the period from 2017 to 2023, as the post-processing of predictions results in more conservative predictions for the initial and final years (2016 and 2024). Over this period, the total area of mining polygons has increased by roughly 24% in the tropics, extending over about 74,100 km² in 2023. In this period, the number of unique mining polygons in covered locations also increased by 23% to slightly more than 19,000 in the year 2023, while their average size remained stagnant.




Figure 3. Summary of the predictions for the **full dataset**, and separated into regions — the **Americas**, **Africa**, and **Asia and Oceania**. The left panel shows the increase in area (in 10,000 km²), the center panel the number of individual mining polygons (in thousands), and the right panel shows the average size of these mining polygons (in km²). The cross (×) indicates values for the ground truth.

These dynamics are similar on the regional level. (Figure S2 in the Supplementary Information further decomposes the regional graphs in Figure 3 to show the three countries within them that have the largest share of total mining area.) The largest absolute increase in mining polygon area in covered sites occurred in Asia and Oceania, with a total increase of 6,200 km², or more than 31%, from 2017 to 2023. In terms of total area, the Americas host the most expansive mining sites in 2023, at over 32,600 km², increasing by 23% since 2017. In Africa, the increase in total mining polygon area amounted to 16%, expanding from 12,900 km² in 2017 to 14,900 km² in 2023.

In Asia and Oceania, Indonesia stands out as the country with the largest mining polygon area. From 2017–2023, we observe considerable growth of Indonesian mining sites at 40%, making it also the country

with the largest absolute increase in mining polygon area globally. Two events that relate to this trend are a Chinese ban of coal imports from Australia in December 2020 and a 2020 Mining Law Amendment that relaxed environmental requirements.³⁴ Notably, the annual mining polygon area in Indonesia is relatively noisy, as the country houses some of the largest single mining sites. Other notable countries in the region are India, which showed stable growth, and Myanmar. There, the area of mining polygons increased until 2020, and began to drop from 2021 onward, coinciding with the ongoing ‘Spring Revolution’.

In the Americas, Brazil is the country with the most expansive mining polygons. Significant growth occurred from 2019–2022, during the administration of Jair Bolsonaro. Most of this expansion took place in and along the fringes of the Amazon rainforest (cf., Figure 1), corroborating early concerns about potential repercussions of legislative changes and increases in anti-environmental rhetoric.³⁵ Chile, which is only partially covered by our dataset, comes a close second. Mining polygons in the country are, on average, larger than in other countries in the region. From 2017 to 2023, the area of mining polygons in the country increased by about 25%.

In Africa, the country with the highest share of mining polygons and area is, despite it being only covered partially, South Africa. The country has been experiencing stagnating mineral production since the early 2010s,³⁶ which is reflected by flat figures for the change of area and the number of mining polygons in our dataset. Ghana drove most of the regional growth, with steady expansion that correlates with rising international gold prices. Gold accounts for around 95% of the country’s mineral revenues, and much of it is produced in small-scale artisanal mines.²⁵ For Africa, Figure 3 shows a substantial increase in the number of mining polygons and moderate area increases, which may be indicative of artisanal mining. However, artisanal mining sites, which are widespread across Africa,³⁷ are underrepresented in our dataset.

Applications

The introduced dataset contributes a time dimension to mining polygons at known locations. This enables various analytical applications and extensions, while maintaining certain limitations that we discuss next. First, it allows for more nuanced analyses of the direct and indirect impacts of mines on their surrounding environments, such as forest loss or vegetation impacts.^{9,14} Second, it facilitates monitoring of mining sites, which can inform re-naturalization projects for abandoned mines. Third, changes in the area of mining sites can serve as a proxy for on-site activity. This enables more accurate assessments of mining-related impacts where data is incomplete or scarce, and facilitates analyses in settings where no conventional data is available, e.g., due to conflict. For example, in analyses of societal impacts like conflict and corruption,⁶ environmental effects such as air and water pollution,¹² and their resulting socioeconomic consequences.

Advantages

Our automated approach offers several distinct advantages. The segmentation process is internally consistent and can be adjusted to meet specific requirements, such as balancing the cost of false positives versus false negatives. The method also enables efficient scalability in both spatial and temporal dimensions. As new satellite images become available, we can easily expand the temporal scope of the dataset. Additionally, we can incorporate more frequent satellite imagery to assess intra-term changes in mining footprints. This could, e.g., help reveal transitory substitution of agricultural activities due to meteorological shocks in subsistence economies. The spatial scope can be extended by incorporating new mining locations — in regions beyond the tropics, or ones that are only sparsely covered by our sources (e.g., in the Congo Basin).

Current limitations

The dataset and approach presented here have noteworthy limitations.

First, polygons of individual mining sites should be used with caution. While the model performs well on average and post-processing reduces the noise in our predictions, the results are not free of errors. This is particularly problematic (a) where mining sites are imprecisely delineated in the ground truth (e.g., artisanal mining sites in Western Africa or the Amazon), (b) in areas where the environment provides limited contrast (e.g., along the western coast of South America). While our predictions compare favorably with the manually delineated datasets used to train the model (see the Methods section), we recommend using any of these datasets at more aggregated levels.

Second, our approach relies on numerous high-resolution satellite images covering vast areas at high frequency. This paper was possible due to the Planet/NICFI program, which provided free satellite imagery of the tropical belt for research purposes. After the program's expiration in January 2025, the reproducibility and extensibility of the presented dataset are impeded. Future work and extensions will require access to other sources of satellite imagery. These are currently prohibitively expensive (in the case of high-resolution sensors) or require sacrifices in terms of the resolution and frequency of available imagery (in the case of Sentinel-2).

Third, our dataset focuses on tracking known mining sites over time, rather than identifying newly established sites far from existing ones. Our source data on mining locations and, consequently, our predictions primarily cover industrial mines and smaller mines in their vicinity, potentially under-representing regions where illegal or informal mining predominates. As a result, our dataset cannot provide a complete overview of the development of tropical mining areas. Enhancing the coverage and delineation of artisanal mining sites, as well as obtaining more comprehensive coverage statistics and validation results, will require additional data.

Methods

In this section, we describe the data and methods used to produce the dataset and our validation strategy.

We begin by describing the satellite imagery used to delineate polygons, as well as the mining polygons that constitute our ground truth.

Satellite imagery

We use high-resolution (< 5 m per pixel) satellite imagery from Planet's PlanetScope constellation of imaging satellites, made available for non-commercial use by NICFI. The imagery covers approximately 45 million square kilometers along the tropical belt, ± 30 degrees of latitude. This includes several important mineral-extracting countries such as Indonesia, the Democratic Republic of the Congo (DRC), and Brazil (see [Figure 1](#)), but excludes, e.g., Australia, the US, Russia, or China, and only partially covers Chile, Argentina, or South Africa.

The dataset spans from December 2015 through December 2024. Imagery is available as biannual composites (December–May, June–November) until May 2020, as a quarterly composite for June–August 2020, and as monthly composites thereafter. Images are provided as so-called ‘quads’, which are compiled from individual daily ‘scenes’ that are captured daily by the constellation of approximately 130 satellites. An illustration of two monthly quads is provided in [Figure S3](#) of the Supplementary Information. This frequent capture rate of imagery results in composite images with relatively high clarity.

For our annual predictions, we use imagery from the second half of the year (June–November). For 2016–2019, we use the provided quads as is, while we select the composites with optimal quality (in terms of minimum cloud coverage) for 2020–2024. While quality metadata only exists for individual scenes

and thus represents worst-case estimates, we observe over 96% clarity on average across our dataset (see [Table S2](#) in the Supplementary Information for details).

Mining polygons

Our study relies on mining polygons from two comprehensive datasets to (a) establish a ground truth to train our model and (b) provide mining locations for prediction.

The first dataset, initially presented in Maus et al. (2020)³⁸ and updated in Maus et al. (2022),²⁶ covers both industrial and artisanal mines. The dataset contains 44,929 mining polygons covering approximately 101,583 km² globally. The original dataset combined undated imagery from Google Satellite, Microsoft Bing Imagery, and Sentinel-2, the updated data is based on Sentinel-2 imagery (10 m per pixel) from 2019. Starting from the SNL Metals and Mining database locations, the authors examined 10 km buffers around known mines to account for expansive operations and nearby (artisanal) mining sites.

The second dataset is an extended version of Tang et al. (2021)³⁹, updated by Tang and Werner (2023),²⁸ and contains 74,548 mining polygons, covering roughly 66,000 km² globally. The authors combined various sources for mine locations, including the SNL database, as well as satellite imagery. Imagery sources include Sentinel-2, Landsat, and high-resolution imagery available in Google Earth Pro (< 1 m per pixel), although no single source was used consistently, and no timestamps are available.

These two datasets have considerable overlap in terms of mine locations but differ in their segmentation approaches. The second dataset²⁸ precisely contours individual mine features, such as waste rock dumps, pits, ponds, tailings dams, and infrastructure. By contrast, Maus et al.²⁶ attempt to cover mining sites in general, often including areas that cannot be strictly classified as mining areas themselves. The Supplementary Information provides examples in [Figure S1](#), showing polygons from both datasets overlaid on Planet/NICFI imagery from 2019 at four locations.

Ground truth

Our ground truth data combines these two state-of-the-art mining polygon datasets, with each presenting distinct challenges and advantages. The primary limitation in both datasets is the lack of exact timestamps. We use 2019 as the reference year, but polygons may have been delineated using earlier satellite imagery. Additional imprecision may stem from heterogeneous satellite imagery, differences in delineation approaches, and human factors, such as individual interpretation styles and fatigue effects.

To maximize the available information, we take the union of both datasets as our ground truth. Experiments using either dataset individually yielded inferior results. Tang and Werner's²⁸ precise delineation proved difficult to reproduce consistently, while using only Maus et al.'s²⁶ dataset produced slightly worse segmentation results than our combined approach.

The union preserves polygons from both sources and maintains outer boundaries where they intersect. The resulting ground truth inherits the broader notion of mining polygons from the first dataset while retaining precise boundary information from the second one. While this approach sacrifices some granularity in distinguishing individual mining features, it creates a more robust foundation for automatic delineation and enables scalable analyses by incorporating the strengths of both source datasets.

Segmentation model

For the delineation of mining areas over time, we use a segmentation model that is based on state-of-the-art transformer architectures from the ML literature.^{40,41} The process of obtaining mining polygons consists of four steps: (1) preprocessing satellite imagery to generate datasets for training and prediction, (2) training the model using our ground truth, (3) predicting mining site delineations for 2016-2024 and

mapping them back to their geospatial locations, and (4) postprocessing the predictions to reduce noise and ensure temporal continuity.

Training the segmentation model requires two sources of data: satellite images and corresponding segmentation masks that indicate mining areas. The model is trained on satellite imagery from 2019, and segmentation masks are only generated for this reference year. The images, however, are also used for prediction and need to be generated for each year and polygon. We generate a square image with twice the maximum side length of the respective polygon to accommodate potential growth over time and downscale them to 512×512 pixels if necessary. The resulting data is split into train/validation/test samples using an 89/10/1 ratio, resulting in approximately 18,300 training samples, 2,100 validation samples, and 200 samples for testing. The test samples were manually selected (from random samples) to ensure they contain only precise, high-quality delineations, avoiding noise in the ground truth.

For the segmentation task, we utilize the SegFormer model.³¹ SegFormer has demonstrated strong performance on benchmark data sets^{42,43} and outperforms older architectures like U-Net,^{44,45} which has previously been used for delineating mining areas in Sub-Saharan Africa.²⁷ The model employs a transformer architecture⁴⁰ that provides a broad receptive field and enables parallel processing.⁴⁶ This allows transformer-based models to efficiently analyze images with global context, leading them to regularly outperform convolutional neural networks like U-Net at segmentation.⁴⁵

We specifically employ SegFormer-B5,³¹ the largest model variant. The encoder undergoes pre-training on the Imagenet-1k dataset,⁴⁷ before fine-tuning the whole model on our mining dataset for about 20 epochs. We train on two NVIDIA A30 GPUs, employing class-balanced loss to address the imbalance between mining and non-mining areas. Additionally, we apply data augmentation techniques, including random horizontal and vertical flipping, random cropping, random resizing, and online hard example mining⁴⁸ to focus on challenging cases.

Due to the class imbalance, we use the *mean Intersection over Union* (mIoU) to assess the model's predictive performance. The model achieves an mIoU of 65.49% and a mean accuracy of 95.67% on the test set when applying an optimized probability threshold $p \geq 0.625$ at the model's probability output. For a selected threshold of $p \geq 0.55$, chosen to match the ground truth in 2019, the model receives an mIoU of 63.85% and a mean accuracy of 93.21%. For a rough comparison, a U-Net model trained with similar parameters achieved an mIoU of 56.69% and a mean accuracy of 93.35%.

Prediction

The prediction process, illustrated in [Figure S4](#) of the Supplement Information, generates mining polygons for all locations for 2016–2024. The model outputs a probability for each pixel that indicates its confidence in the presence of mining. To classify a certain pixel as part of a mine, we use a probability threshold that can be adjusted to tweak, e.g., the false positive and false negative rates of the predictions. Lower thresholds result in more and larger predicted mining polygons, and thus higher recall, while higher thresholds lead to fewer but more precise predictions.

To illustrate the flexibility that this parameter affords, we follow two strategies in setting it. First, we consider multiple thresholds and manually optimize to $p \geq 0.55$ to emulate the ground truth in the 'Reference validation' section below. This facilitates comparison and highlights the model's ability to interpolate between high-recall and high-precision approaches. As an alternative, we set the parameter by optimizing the mIoU over the validation set, reaching an mIoU of 65.49%. The corresponding threshold of $p \geq 0.625$ means that only high-confidence pixels are classified as mines. The resulting dataset exceeds the precision and specificity of both the ground truth and its individual components in segmenting mines but sacrifices recall.

To ease comparison with our source datasets, we use $p \geq 0.55$ as the default threshold, where not

mentioned otherwise. After creating the binary pixel-level predictions, we apply a contour-finding algorithm⁴⁹ to transform these segmentation masks to polygons and map the polygons back to their geospatial coordinates using the previously calculated bounding boxes.

We implement postprocessing to ensure temporal continuity and to reduce noise in the predictions. We remove any polygons that do not have an intersecting polygon in the previous *or* subsequent year. This allows us to borrow information across years and reflects the assumption that mines (and their corresponding polygons) do not emerge and disappear within a single year. Hence, the years 2016 and 2024, which only have a single ‘neighboring’ year to compare to, feature a lower *number* of polygons, while the impact on mining areas is limited.

To investigate potential bias from this postprocessing step, we experimented with applying a 100 m buffer to the predicted polygons before the intersection, assuming that new mines may emerge or disappear near existing ones. We found that this buffer had little impact on the result. This suggests that the postprocessing step primarily removes isolated polygons, which may represent false positives and can be considered as noise.

Lastly, we merge multiple, intersecting predictions that may stem from locations with overlapping bounding boxes and apply hole-filling to remove fragments inside the predictions.

Reference validation

To assess our model in relation to existing datasets in the literature and fine-tune probability thresholds for comparability, we use a set of labeled validation points. We rely on the set of validation points from Maus et al.,²⁶ which were generated within a 10 km buffer around the mines. Of their 1,220 original validation points, 450 are covered by Planet/NICFI. We consider a random subset of 200 points and manually label them as either mine or non-mine for each year using Planet/NICFI imagery. Based on these points, we can systematically compare (1) our model’s predictions, (2) the ground truth, and (3) its two constituting datasets.

For reference, Tang and Werner²⁸ use 200 validation polygons that are not readily available and have limited coverage in our area of interest.

Before assessing the results, we note two drawbacks related to the sampling strategy used to generate the validation points. The locations are not sampled randomly but stratified across the mine and non-mine categories, as segmented by Maus et al.²⁶ As a result, (likely) non-mine labels are spread out across their (large) area of interest, while mine labels from within their polygons are overrepresented. Inherent to this strategy, the labels offer almost no meaningful variation over time. Points from the vicinity of mining polygons, which may develop into mining areas over time, are undersampled, and 181 of the 200 points do not change their label over time. We address this by reporting averages of the yearly validation results. Nevertheless, the validation exercise must be seen in the context of the Maus et al.²⁶ dataset, from which it was sampled.

Our predictions achieve an average accuracy of 87.7% and an average Matthews Correlation Coefficient (MCC) of 71.7% for a probability threshold of $p \geq 0.55$ on the validation points, indicating performance that is comparable to the source data. For the $p \geq 0.625$ threshold, the average accuracy is 83.3%, and the average MCC is 61.5%. By comparison, the ground truth achieves an average accuracy of 88.1% and MCC of 77.5%; the polygons of Maus et al.²⁶ achieve values of 88.6% and 78.3%, respectively, although these values benefit from the stratified sampling of validation points. Meanwhile, the polygons of Tang and Werner²⁸ achieve an average accuracy of 86.8% and MCC of 69.7%.

Further validation results, including ones for a range of probability thresholds, are visualized in Figure S5 of the Supplementary Information. The scores of our selected threshold of $p \geq 0.55$ lie between the two sources for our ground truth, demonstrating how our method can combine their liberal and precise

delineation styles. This indicates the model's ability to effectively integrate information from both sources, identify a shared definition of mining polygons, and accurately extrapolate it to nine different timestamps.

What is notable is that our model can go beyond the performance of either dataset used to train it. Due to the imbalance between non-mine and mine areas, we particularly value precision (user's accuracy, UA) in segmenting mines over their recall (producer's accuracy, PA), although the opposite would also be possible. The selected threshold of $p \geq 0.55$ already achieves an average precision of 84.1%, exceeding the ground truth (73.5%) and its components (74.4% and 82.3%), while the higher threshold of $p \geq 0.625$ even reaches 90.3%. As a result, predictions feature considerably lower false positive rates at 7.1% and 3.0% than the ground truth (17.5%), Maus et al.²⁶ (16.8%), or Tang and Werner²⁸ (8.0%). This highlights the adaptability of our modeling approach, which can be fine-tuned to the task at hand.

Data availability

The dataset of mining polygons is available in multiple formats to ensure accessibility. An interactive map that allows the visualization of yearly predictions against Sentinel-2 satellite imagery is available at [layers.at](#). The full dataset, which includes the hand-validated test set and validation points, is made available on Zenodo under [10.5281/zenodo.17034288](#). The training and validation datasets are provided on [kaggle.com](#). The dataset is licensed under the Open Database License (ODbL, available at [opendatacommons.org](#)). Any rights in individual contents of the database are licensed under the Database Contents License (DbCL, available at [opendatacommons.org](#)). The dataset and validation points are provided in the GeoPackage (GPKG) format and contain information on the polygons themselves, the countries where they are located, the years in which they are segmented, and the areas of polygons.

Code availability

All code used to produce the dataset and results of this paper is available under the GNU General Public License v3.0 (GPLv3, available at [gnu.org](#)) at [github.com](#). The repository includes comprehensive documentation and usage instructions. Scripts were written in Python and R, with geospatial processing heavily utilizing the 'GeoPandas',⁵⁰ 'Shapely',⁵¹ and 'sf'⁵² packages. The model was implemented using 'PyTorch'⁵³ and 'MMSegmentation'.⁵⁴

References

1. Shukla, P. R. *et al.* *Climate Change 2022: Mitigation of Climate Change* (Cambridge University Press, Cambridge, 2022).
2. Luckeneder, S., Giljum, S., Schaffartzik, A., Maus, V. & Tost, M. Surge in global metal mining threatens vulnerable ecosystems. *Glob. Environ. Chang.* **69**, 102303, [10.1016/j.gloenvcha.2021.102303](#) (2021).
3. International Energy Agency (IEA). Energy technology perspectives 2023. Tech. Rep., IEA (2023).
4. Hund, K., La Porta, D., Fabregas, T. P., Laing, T. & Drexhage, J. Minerals for climate action: The mineral intensity of the clean energy transition. Tech. Rep., World Bank Group (2023).
5. Kossoff, D. *et al.* Mine tailings dams: Characteristics, failure, environmental impacts, and remediation. *Appl. Geochem.* **51**, 229–245, [10.1016/j.apgeochem.2014.09.010](#) (2014).
6. Berman, N., Couttenier, M., Rohner, D. & Thoenig, M. This mine is mine! How minerals fuel conflicts in Africa. *Am. Econ. Rev.* **107**, 1564–1610, [10.1257/aer.20150774](#) (2017).

7. Sonter, L. J., Dade, M. C., Watson, J. E. M. & Valenta, R. K. Renewable energy production will exacerbate mining threats to biodiversity. *Nat. Commun.* **11**, 1–6, [10.1038/s41467-020-17928-5](https://doi.org/10.1038/s41467-020-17928-5) (2020).
8. Liu, Y. *et al.* A review of water pollution arising from agriculture and mining activities in central asia: Facts, causes and effects. *Environ. Pollut.* **291**, 118209, [10.1016/j.envpol.2021.118209](https://doi.org/10.1016/j.envpol.2021.118209) (2021).
9. Giljum, S. *et al.* A pantropical assessment of deforestation caused by industrial mining. *Proc. Natl. Acad. Sci.* **119**, e2118273119, [10.1073/pnas.2118273119](https://doi.org/10.1073/pnas.2118273119) (2022).
10. Villén-Pérez, S., Anaya-Valenzuela, L., Conrado da Cruz, D. & Fearnside, P. M. Mining threatens isolated indigenous peoples in the Brazilian Amazon. *Glob. Environ. Chang.* **72**, 102398, [10.1016/j.gloenvcha.2021.102398](https://doi.org/10.1016/j.gloenvcha.2021.102398) (2022).
11. Berman, N., Couttenier, M. & Girard, V. Mineral resources and the salience of ethnic identities. *Econ. J.* **133**, 1705–1737, [10.1093/ej/uead018](https://doi.org/10.1093/ej/uead018) (2023).
12. Macklin, M. G. *et al.* Impacts of metal mining on river systems: A global assessment. *Science* **381**, 1345–1350, [10.1126/science.adg6704](https://doi.org/10.1126/science.adg6704) (2023).
13. Mwelwa, S., Chungu, D., Tailoka, F., Beesigamukama, D. & Tanga, C. Biotransfer of heavy metals along the soil-plant-edible insect-human food chain in Africa. *Sci. The Total. Environ.* **881**, 163150, [10.1016/j.scitotenv.2023.163150](https://doi.org/10.1016/j.scitotenv.2023.163150) (2023).
14. Vashold, L., Pirich, G., Heinze, M. & Kuschnig, N. Mines–rivers–yields: Downstream mining impacts on agriculture in Africa. *WU Work. Pap.* [10.57938/7598a00e-3d24-4691-823e-aeb76032ae44](https://doi.org/10.57938/7598a00e-3d24-4691-823e-aeb76032ae44) (2024).
15. Aragón, F. M. & Rud, J. P. Polluting industries and agricultural productivity: Evidence from mining in Ghana. *The Econ. J.* **126**, 19802011, [10.1111/eco.12244](https://doi.org/10.1111/eco.12244) (2015).
16. von der Goltz, J. & Barnwal, P. Mines: The local wealth and health effects of mineral mining in developing countries. *J. Dev. Econ.* **139**, 1–16, [10.1016/j.jdeveco.2018.05.005](https://doi.org/10.1016/j.jdeveco.2018.05.005) (2019).
17. Bazillier, R. & Girard, V. The gold digger and the machine. evidence on the distributive effect of the artisanal and industrial gold rushes in Burkina Faso. *J. Dev. Econ.* **143**, 102411, [10.1016/j.jdeveco.2019.102411](https://doi.org/10.1016/j.jdeveco.2019.102411) (2020).
18. Ofosu, G., Dittmann, A., Sarpong, D. & Botchie, D. Socio-economic and environmental implications of artisanal and small-scale mining (asm) on agriculture and livelihoods. *Environ. Sci. & Policy* **106**, 210–220, [10.1016/j.envsci.2020.02.005](https://doi.org/10.1016/j.envsci.2020.02.005) (2020).
19. Bebbington, A., Hinojosa, L., Bebbington, D. H., Burneo, M. L. & Warnaars, X. Contention and ambiguity: Mining and the possibilities of development. *Dev. Chang.* **39**, 887–914, [10.1111/j.1467-7660.2008.00517.x](https://doi.org/10.1111/j.1467-7660.2008.00517.x) (2008).
20. van der Ploeg, F. Natural resources: Curse or blessing? *J. Econ. Lit.* **49**, 366–420, [10.1257/jel.49.2.366](https://doi.org/10.1257/jel.49.2.366) (2011).
21. Venables, A. J. Using Natural Resources for Development: Why Has It Proven So Difficult? *J. Econ. Perspectives* **30**, 161–84, [10.1257/jep.30.1.161](https://doi.org/10.1257/jep.30.1.161) (2016).
22. Badeeb, R. A., Lean, H. H. & Clark, J. The evolution of the natural resource curse thesis: A critical literature survey. *Resour. Policy* **51**, 123–134, [10.1016/j.resourpol.2016.10.015](https://doi.org/10.1016/j.resourpol.2016.10.015) (2017).
23. Franks, D. M., Keenan, J. & Hailu, D. Mineral security essential to achieving the Sustainable Development Goals. *Nat. Sustain.* **6**, 21–27, [10.1038/s41893-022-00967-9](https://doi.org/10.1038/s41893-022-00967-9) (2023).

24. Ali, S. H. *et al.* Mineral supply for sustainable development requires resource governance. *Nature* **543**, 367–372, [10.1038/nature21359](https://doi.org/10.1038/nature21359) (2017).
25. Barenblitt, A. *et al.* The large footprint of small-scale artisanal gold mining in Ghana. *Sci. The Total. Environ.* **781**, 146644, [10.1016/j.scitotenv.2021.146644](https://doi.org/10.1016/j.scitotenv.2021.146644) (2021).
26. Maus, V. *et al.* An update on global mining land use. *Sci. data* **9**, 1–11, [10.1038/s41597-022-01547-4](https://doi.org/10.1038/s41597-022-01547-4) (2022).
27. Provenzano, S. & Bull, H. The local economic impact of mineral mining in africa: Evidence from four decades of satellite imagery. *ArXiv e-prints* [10.48550/arXiv.2111.05783](https://arxiv.org/abs/2111.05783) (2021).
28. Tang, L. & Werner, T. T. Global mining footprint mapped from high-resolution satellite imagery. *Commun. Earth & Environ.* **4**, 134, [10.1038/s43247-023-00805-6](https://doi.org/10.1038/s43247-023-00805-6) (2023).
29. Quash, Y., Kross, A. & Jaeger, J. A. G. Assessing the impact of gold mining on forest cover in the Surinamese Amazon from 1997 to 2019: A semi-automated satellite-based approach. *Ecol. Informatics* **80**, 102442, [10.1016/j.ecoinf.2023.102442](https://doi.org/10.1016/j.ecoinf.2023.102442) (2024).
30. Maus, V. & Werner, T. T. Impacts for half of the world’s mining areas are undocumented. *Nature* **625**, 26–29, [10.1038/d41586-023-04090-3](https://doi.org/10.1038/d41586-023-04090-3) (2024).
31. Xie, E. *et al.* Segformer: Simple and efficient design for semantic segmentation with transformers. *Adv. Neural Inf. Process. Syst.* **34**, 12077–12090, [10.48550/arXiv.2105.15203](https://arxiv.org/abs/2105.15203) (2021).
32. Archi Indonesia. Public expose 2024. https://www.idx.co.id/StaticData/NewsAndAnnouncement/ANNOUNCEMENTSTOCK/From_EREP/202406/fdfde002ad_a4aef307c1.pdf (2024).
33. Verweijen, J., Schouten, P., Simpson, F. O. & Pascal, C. Z. Conservation, conflict and semi-industrial mining: the case of eastern DRC. *Analysis Policy Briefs* (2022).
34. Werner, T. T. *et al.* Patterns of infringement, risk, and impact driven by coal mining permits in Indonesia. *Ambio* **53**, 242–256, [10.1007/s13280-023-01944-y](https://doi.org/10.1007/s13280-023-01944-y) (2024).
35. Ferrante, L. & Fearnside, P. M. Brazil threatens indigenous lands. *Science* **368**, 481–482, [10.1126/science.abb6327](https://doi.org/10.1126/science.abb6327) (2020).
36. Statistics South Africa. P2401 mining: Production and sales (2024).
37. Hilson, G. & McQuilken, J. Four decades of support for artisanal and small-scale mining in sub-Saharan Africa: A critical review. *Extr. Ind. Soc.* **1**, 104–118, [10.1016/j.exis.2014.01.002](https://doi.org/10.1016/j.exis.2014.01.002) (2014).
38. Maus, V. *et al.* A global-scale data set of mining areas. *Sci. Data* **7**, 1–13, [10.1038/s41597-020-00624-w](https://doi.org/10.1038/s41597-020-00624-w) (2020).
39. Tang, L., Werner, T. T., Heping, X., Jingsong, Y. & Zeming, S. A global-scale spatial assessment and geodatabase of mine areas. *Glob. Planet. Chang.* **204**, 103578, [10.1016/j.gloplacha.2021.103578](https://doi.org/10.1016/j.gloplacha.2021.103578) (2021).
40. Vaswani, A. *et al.* Attention is all you need. *Conf. on Neural Inf. Process. Syst.* **30**, [10.48550/arXiv.1706.03762](https://arxiv.org/abs/1706.03762) (2017).
41. Li, X. *et al.* Transformer-based visual segmentation: A survey. *ArXiv e-prints* [10.48550/arXiv.2304.09854](https://arxiv.org/abs/2304.09854) (2023).
42. Papers with Code. Semantic segmentation on ADE20K. <https://paperswithcode.com/sota/semantic-segmentation-on-ade20k> (2025).

- 43.** Papers with Code. Semantic segmentation on Cityscapes. <https://paperswithcode.com/sota/semantic-segmentation-on-cityscapes> (2025).
- 44.** Ronneberger, O., Fischer, P. & Brox, T. U-Net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, 234–241, [10.1007/978-3-319-24574-4_28](https://doi.org/10.1007/978-3-319-24574-4_28) (Springer, Cham, Switzerland, 2015).
- 45.** Sourget, T., Hasany, S. N., Mériadeau, F. & Petitjean, C. Can SegFormer be a true competitor to U-Net for medical image segmentation? In *Medical Image Understanding and Analysis*, 111–118, [10.1007/978-3-031-48593-0_8](https://doi.org/10.1007/978-3-031-48593-0_8) (Springer, Cham, Switzerland, 2023).
- 46.** Luo, W., Li, Y., Urtasun, R. & Zemel, R. Understanding the effective receptive field in deep convolutional neural networks. In *Conference on Neural Information Processing Systems*, 4905–4913, [10.48550/arXiv.1701.04128](https://doi.org/10.48550/arXiv.1701.04128) (Curran Associates Inc., 2016).
- 47.** Russakovsky, O. *et al.* ImageNet large scale visual recognition challenge. *Int. J. Comput. Vis. (IJCV)* **115**, 211–252, [10.1007/s11263-015-0816-y](https://doi.org/10.1007/s11263-015-0816-y) (2015).
- 48.** Shrivastava, A., Gupta, A. & Girshick, R. Training region-based object detectors with online hard example mining. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 27–30, [10.1109/CVPR.2016.89](https://doi.org/10.1109/CVPR.2016.89) (IEEE, 2016).
- 49.** Suzuki, S. & Be, K. Topological structural analysis of digitized binary images by border following. *Comput. Vision, Graph. Image Process.* **30**, 32–46, [10.1016/0734-189X\(85\)90016-7](https://doi.org/10.1016/0734-189X(85)90016-7) (1985).
- 50.** Jordahl, K. *et al.* Geopandas 1.0.1, [10.5281/zenodo.2585848](https://doi.org/10.5281/zenodo.2585848) (2024).
- 51.** Gillies, S. & contributors, S. Shapely: A Python package for manipulation and analysis of planar geometric objects. <https://shapely.readthedocs.io/en/stable/index.html> (2024).
- 52.** Pebesma, E. & Bivand, R. *Spatial Data Science: With applications in R* (Chapman and Hall/CRC, 2023).
- 53.** Paszke, A. *et al.* PyTorch: An imperative style, high-performance deep learning library. In *Conference on Neural Information Processing Systems*, 8026–8037, [10.48550/arXiv.1912.01703](https://doi.org/10.48550/arXiv.1912.01703) (Curran Associates Inc., 2019).
- 54.** MMSegmentation. OpenMMLab Semantic Segmentation Toolbox and Benchmark. <https://github.com/open-mmlab/mmsegmentation> (2020).

Acknowledgements

The authors gratefully acknowledge financial support from the Oesterreichische Nationalbank (OeNB anniversary fund, project No. 18799; NK, LV, PS) and the City of Vienna (Hochschuljubiläumsfonds, project No. H-457973/2023; LV, PS). Data access under the NICFI Satellite Data Program enabled this research. Computations were performed on the WUCluster, a facility supported by the Core Services of the Vienna University of Economics and Business.

Author contributions

N.K and P.S conceived the study. P.S implemented the method. All authors wrote and reviewed the manuscript.

Competing interests

The authors declare no competing interests.

Supplementary Information

Here we provide some additional figures and tables that allow for more detailed insights into the produced dataset as well as the underlying ground truth data and satellite imagery.

[Figure S1](#) shows satellite imagery from 2019 for four mine sites across the covered area, overlaid with polygons from the two datasets that constitute the ground truth used for training the segmentation model. It indicates that the broadly defined delineations of Maus et al. (2022)²⁶ often encompass the more precise delineations by Tang and Werner (2023)²⁸, which also sometimes pick up additional mining features.

[Table S1](#) provides information on the total mining polygon area for countries with a large mining footprint (defined as exceeding 100 km² at least once during the observation period). It extends [Table 1](#) in the manuscript and confirms the general trend of increasing total mining polygon area over the years also within most countries.

[Figure S2](#) decomposes the global version [Figure 3](#) in the manuscript to depict regional summaries of the predictions over time. It shows the development of the total area and number of mining polygons for the three countries within each region that had the highest share of overall mining polygon area separately as well as for the remaining countries aggregated.

[Figure S3](#) and [Table S2](#) provide some additional information about the results from the steps undertaken to select the images with the highest clarity (i.e., lowest cloud and heavy haze coverage) from a 6-month window (June–November) for each mining site individually. [Table S2](#) shows that the clarity of selected yearly satellite images for the period 2020–2024 is on average comparable to or even exceeds the one for the period 2016–2019, where we take the biannual composites as is. [Figure S3](#) provides an example of the importance of this procedure to select highest-clarity images, showing two composite images from October and December 2024 for the area surrounding the Toka Tindung mining site.

[Figure S4](#) illustrates the prediction process while [Figure S5](#) shows the validation results for our predictions. [Figure S6](#) shows some validation metrics discussed in the Methods section of the paper over time for two different probability thresholds. It shows limited variability in most of these metrics over time, which mostly stems from a lack of variation in the labels of the validation points. Yet, it can be discerned that our modeling approach achieves performance comparable to the ground truth and its constituting datasets over the year. Differences in metrics across the selected probability thresholds emulate the ones shown in [Figure S5](#).




Figure S1. Four examples of the source polygons for our ground truth.^{26,28} The mine sites are located in Indonesia (Toka Tindung, from Figure 2), the DRC (Kapulo Copper Mine), Chile (Minera Sierra Gorda), and Suriname (south of the Brokopondo Reservoir) against 2019 Planet/NICFI imagery. Note the different zoom levels of the top and bottom rows. Imagery ©2019 Planet Labs Inc.

Country	Total Mining Polygon Area (km ²)								
	2016	2017	2018	2019	2020	2021	2022	2023	2024
Indonesia	13,266	12,425	12,376	13,761	15,281	16,859	15,746	17,497	17,801
Brazil	7,829	8,785	8,002	8,023	8,548	8,215	9,764	9,342	9,360
Chile	6,293	5,748	5,746	7,161	7,269	6,434	7,561	7,176	7,218
South Africa	7,469	6,624	5,901	6,706	6,324	6,044	6,434	6,459	6,558
India	3,502	3,800	3,590	4,140	4,136	3,959	4,322	4,367	4,114
Myanmar	2,670	2,497	2,667	2,607	3,800	3,129	3,099	2,769	2,890
Bolivia	557	787	806	1,025	910	1,103	1,108	2,035	2,695
Guyana	2,082	2,245	2,368	2,246	2,380	3,213	2,942	3,197	2,570
Suriname	1,906	1,824	1,902	1,986	1,976	2,111	2,119	2,382	2,430
Peru	2,328	2,502	2,860	2,735	2,816	3,287	4,286	3,172	2,165
Venezuela	1,183	1,354	1,285	1,341	1,414	1,442	1,310	1,528	1,688
Ghana	1,143	735	957	1,189	979	1,158	1,411	1,516	1,641
Namibia	596	1,037	925	1,217	1,135	1,043	1,503	1,381	1,246
Colombia	961	908	885	1,013	927	1,067	1,083	1,117	1,083
Mexico	812	847	864	957	931	959	1,060	1,023	998
Argentina	816	916	898	933	931	884	816	875	963
Zambia	732	668	660	677	803	745	757	814	830
Dem. Rep. Congo	615	622	629	680	732	692	732	794	806
Zimbabwe	505	458	484	472	508	585	607	697	674
Angola	500	440	421	442	487	540	479	543	562
Vietnam	478	458	443	446	419	519	536	516	523
Philippines	372	370	407	415	395	428	456	489	430
Botswana	297	343	366	353	351	354	334	379	360
Malaysia	268	273	270	257	237	300	320	331	331
Mali	252	219	232	276	251	295	358	355	317
Guinea	214	239	213	263	242	290	278	294	236
France	193	227	257	218	189	216	247	242	233
Burkina Faso	188	223	190	206	184	191	203	204	227
Thailand	200	182	196	224	204	231	210	198	221
Mozambique	188	177	197	120	211	224	210	219	218
Tanzania	212	199	171	208	190	212	204	236	212
Papua New Guinea	150	197	219	220	240	407	221	218	199
Cuba	147	124	138	137	133	166	180	172	158
Senegal	104	127	143	128	134	133	130	127	149
Niger	142	181	171	178	195	138	207	138	143
Sierra Leone	93	113	123	132	143	134	139	152	136
Ecuador	62	70	74	92	118	73	117	119	121
Côte d'Ivoire	60	74	96	99	100	113	129	132	111

Table S1. Total mining polygon area in the countries with a large mining footprint ($\geq 100 \text{ km}^2$) in the dataset over time.




Figure S2. Regional summaries of the predictions over time: the left column shows total mining areas, and the right the number of mine polygons. Note, in particular, the strong increases in area for Chile and Peru in 2022 (top row), the sharp decrease in area following years of steady expansion for Indonesia in 2022 (middle row), and the steady expansion in the number of mining polygons coupled with only modest increases in area for Ghana (bottom row).




Figure S3. Two satellite images (Planet/NICFI) of the Toka Tindung mine (from Figure 2) for 2024. The left image is a composite from October 2024 and features high clarity (99.9%), while the latest image from December 2024 has clearly visible artifacts, reducing its clarity. Our approach selects the highest-clarity image from a six-month window, leading to high-quality images throughout. Imagery ©2024 Planet Labs Inc.

Year	Mean	SD	Percentiles				
			1st	5th	10th	50th	90th
2016	94.29	(12.32)	38.00	67.00	81.36	99.64	100
2017	95.24	(10.21)	46.94	72.83	87.00	99.96	100
2018	94.55	(10.13)	57.19	67.54	80.52	99.33	100
2019	97.05	(8.15)	63.43	73.79	89.48	100.0	100
2020	97.44	(7.28)	65.46	77.64	91.00	100.0	100
2021	97.48	(7.57)	62.73	78.07	94.25	100.0	100
2022	95.61	(11.25)	53.70	60.80	84.11	100.0	100
2023	95.80	(10.77)	53.74	65.16	82.20	100.0	100
2024	96.01	(8.97)	66.79	72.14	79.16	100.0	100

Table S2. Summary statistics for the visibility of selected yearly satellite images. Note that the visibility metadata is only available on a per-scene basis; we average scene-level characteristics to compute the relevant quad-level information.




Figure S4. Prediction process for three mining sites in Brazil (near Peixoto de Azevedo), Mozambique (Balama graphite mine), and Indonesia (Toka Tindung, see Figure 2) over satellite imagery from 2019. The left column shows ground truth polygons, the center column shows the model's probabilistic prediction, and the right column shows the predicted polygons for thresholds of $p \geq 0.55$ and $p \geq 0.625$. Imagery ©2019 Planet Labs Inc.




Figure S5. Reference validation results for our predictions (dots), the ground truth (solid line), and its components (dashed and dotted lines). Our predictions are evaluated at different probability thresholds for segmentation (along the horizontal axis). Two selected thresholds are highlighted: $p \geq 0.55$, the default value due to its comparability to the ground truth, and $p \geq 0.625$, which achieves the optimal mIoU. Due to the lack of variability over time, the values are averaged over the years; a variant of the plot with year-specific results is provided in [Figure S6](#) in the Supplementary Information.




Figure S6. Reference point validation for individual years.